

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
import time
from scipy.optimize import minimize
```

```
In [2]: def timeit(f):

    def timed(*args, **kw):

        ts = time.time()
        result = f(*args, **kw)
        te = time.time()

        print('func:%r took: %2.4f sec' % (f.__name__, te-ts))
        return result

    return timed
```


Question 1: Bisection VS Golden Section

1. a. Placing e in the bisector of the larger interval $[a, b]$ is better than placing it in $[b, d]$ because it reduces the search space by a larger amount, thus making it a more effective reduction.
- b. The new interval is $[a, e, b]$ and the search space is reduced by 0.25.
- e. Step 2: $13/32$ or 0.4063 Step 3: $33/128$ or 0.2578

1d


```
In [3]: from IPython.display import Image
Image(filename='image1.png')
```

Out [3]:



```
In [4]: from IPython.display import Image
Image(filename='image2.png')
```

Out [4]:



Are my trees ugly? Yes. But they were made with love <3

1f

Golden section converges to a minimum (or maximum) faster than bisection does. In theory this makes sense, but when I typed out a table to compare I saw that bisection (best-case scenario) does better than golden section. I guess golden section is still better because any other case of bisection does not converge as fast, and realistically you will probably rarely get best-case scenario when performing bisection.

```
In [5]: from IPython.display import Image
        Image(filename='image3.png')
```

Out [5]:

A	B	C	D	E
Number of steps	Golden section	Bisection (best case)	Bisection (ok case)	Bisection (worst case)
1	0.618	0.5	0.5	0.75
2	0.382	0.25	0.375	0.5
3	0.236	0.125	0.25	0.375

Question 2: Steepest Descent

```
In [6]: def function3d(point):
        """
        This function calculates the value of a function of x and y at a certain point.
        Parameters
        -----
        x, y: the point to evaluate the function at

        Return
        -----
        the function value at the point
        """
        x, y = point

        return x**4 - x**2 + y**2 + 2*x*y - 2
```

```
In [7]: def first_derivative(point):
        """
        This function calculates the first derivative of a function at a given point.
        Parameters
        -----
        point: list. Not really a list, more like comma separated variables
        A point, x,y

        Returns
        -----
        first_derivative: np.array
        The first derivative (gradient) at the point entered
        """
        x, y = point
        dfdx = (4 * x**3) - (2*x) + 2*y
        dfdy = (2*y) + 2*x

        return np.array([dfdx, dfdy])
```

2a

```
In [8]: @timeit

def steepest_descent_one_step(func, gradient, x0, alpha=0.1):
    """
    This function performs a single step of the steepest descent algorithm.
    Parameters
    -----
    func:
    function to optimize

    gradient: np.array
    gradient of a function

    x0: np.array
    starting point
```

```

num_iter: int
number of times the for loop will run

alpha: float
stepsize

Returns
-----
x1 : new position
value: function evaluated at that x1
"""
point_searched = []
# calculate the initial function value
prev_value = func(x0)

# print search progress and keep track of the points searched
print(f"searching at {x0} with function value {prev_value}")
point_searched.append(x0)

# calculate the gradient at the current point
g = gradient(x0)

# calculate next point and its function value
x1 = x0 - g * alpha
#value = func(x1)

return print(f"New position is {x1}")

```

In [9]: `steepest_descent_one_step(function3d, first_derivative, np.array([1.5, 1.5]))`

```

searching at [1.5 1.5] with function value 7.5625
New position is [0.15 0.9 ]
func:'steepest_descent_one_step' took: 0.0006 sec

```

This is a good optimization step because it moves towards the minimum. We will increase the stepsize * 1.2 in the next step.

2b

```

In [10]: from pylab import *
import numpy.linalg as LA

@timeit
def steepest_descent(func, func_gradient, x0, alpha, tol):
    """
    This function finds a local minimum using the steepest descent method.
    Parameters
    -----
    func:
    The function whose minima we plan to find (inputted as a function)

    func_gradient:
    first_derivative of func
    """

```

```

x0: np.array
position from which we start searching for the minima

alpha: float
step size

tol: float
tolerance value

Returns
-----
Starting point: entered x0
Evaluation: value of function at x0
Path to minimum: lists all points evaluated on the way to minimum
Steps to converge: counts number of steps until local minimum is reached
"""

count = 0
visited = [x0]
deriv = func_gradient(x0)

while LA.norm(deriv) > tol and count < 1e6:
    # calculate new point position
    x1 = x0 - deriv * alpha

    if func(x1) < func(x0):
        # Check if new value is less than previous value. If so, this is
        # and start the next search step from the current value
        alpha = alpha * 1.2
        x0 = x1
        deriv = func_gradient(x0)
        visited.append(x1)

    else:
        # If new_value > prev_value, we are moving in the wrong direction
        # and redo that step starting from the previous value.
        alpha = alpha * 0.5
        #visited.append(x0)

    count += 1
return {"Starting point": x0, "Evaluation": func(x0), "Path to minimum": np.array(visited)}

```

```
In [11]: print(steepest_descent(function3d, first_derivative, np.array([-1.5, 1.5]),
```

```

func:'steepest_descent' took: 0.0005 sec
{'Starting point': array([-1.00000072,  1.0000014 ]), 'Evaluation': -2.99999
9999997453, 'Path to minimum': array([[ -1.5          ,  1.5          ],
    [-0.75         ,  1.5          ],
    [-1.0875        ,  1.32         ],
    [-1.04004413    ,  1.25304        ],
    [-1.05492903    ,  1.17942863    ],
    [-1.00779561    ,  1.12779615    ],
    [-1.05181525    ,  1.0680762     ],
    [-0.98988976    ,  1.06322071    ],
    [-1.03043727    ,  1.03694491    ],
    [-1.00445425    ,  1.03554582    ],
    [-1.00784819    ,  1.02752454    ],
    [-1.00410623    ,  1.021433       ],
    [-1.00440363    ,  1.01499603    ],
    [-1.00122119    ,  1.01027389    ],
    [-1.00344611    ,  1.005431       ],
    [-0.99963579    ,  1.00479389    ],
    [-1.00218339    ,  1.00280712    ],
    [-1.00030255    ,  1.00266297    ],
    [-1.00062139    ,  1.00200836    ],
    [-1.00025502    ,  1.00154679    ],
    [-1.00036336    ,  1.00103092    ],
    [-0.99998637    ,  1.00071101    ],
    [-1.0002104     ,  1.00050266    ],
    [-1.00002076    ,  1.00040182    ],
    [-1.00014415    ,  1.00024404    ],
    [-1.00002569    ,  1.00021923    ],
    [-1.00005275    ,  1.00016153    ],
    [-1.00001618    ,  1.00012262    ],
    [-1.00003409    ,  1.00007692    ],
    [-0.99998591    ,  1.00005486    ],
    [-1.00002464    ,  1.00003355    ],
    [-0.99999138    ,  1.00003024    ],
    [-1.0000077     ,  1.0000216     ],
    [-1.00000319    ,  1.00001789    ],
    [-1.00000381    ,  1.00001317    ],
    [-1.00000155    ,  1.00000957    ],
    [-1.00000239    ,  1.00000587    ],
    [-0.99999901    ,  1.00000395    ],
    [-1.00000196    ,  1.00000231    ],
    [-0.99999897    ,  1.00000217    ],
    [-1.00000072    ,  1.0000014 ]]), 'Steps to converge': 41}

```

This function took 41 good steps to converge. The total number of steps taken was 51.

2c

```

In [12]: # Conjugate gradient
x0 = np.array([-1.5, 1.5])

CG_method = minimize(function3d, x0, method='CG', options={'gtol': 1e-5, 'di
print(CG_method)

```

```

Optimization terminated successfully.
    Current function value: -3.000000
    Iterations: 6
    Function evaluations: 39
    Gradient evaluations: 13
message: Optimization terminated successfully.
success: True
status: 0
    fun: -2.9999999999997273
     x: [-1.000e+00  1.000e+00]
    nit: 6
   jac: [ 2.414e-06  5.364e-07]
  nfev: 39
  njev: 13

```

```

In [13]: # BFGS
x0 = np.array([-1.5, 1.5])

BFGS_method = minimize(function3d, x0, method='BFGS', options={'gtol': 1e-5},
print(BFGS_method)

```

```

Optimization terminated successfully.
    Current function value: -3.000000
    Iterations: 7
    Function evaluations: 27
    Gradient evaluations: 9
message: Optimization terminated successfully.
success: True
status: 0
    fun: -2.999999999999986
     x: [-1.000e+00  1.000e+00]
    nit: 7
   jac: [ 4.172e-07  2.384e-07]
 hess_inv: [[ 1.244e-01 -1.270e-01]
             [-1.270e-01  6.174e-01]]
  nfev: 27
  njev: 9

```

In terms of number of steps, both conjugate gradient (CG) and BFGS are more efficient than steepest descent. CG is only a bit more efficient (39 steps instead of 41) while BFGS takes only 27 steps.

Question 3: Local optimization and machine learning using Stochastic Gradient Descent (SGD)

```

In [14]: def rosenbrock_function3d(point):
        """
        This function calculates the value of a function of x and y at a certain
        Parameters
        -----
        x, y: the point to evaluate the function at

        Return
        -----

```



```

    the function value at the point
    """
    x, y = point

    return (1 - x)**2 + 10 * (y - (x**2))**2

```

```

In [15]: def rosenbrock_gradient(point):
    """
    This function calculates the first derivative of a function at a given point.
    Parameters
    -----
    point: list. Not really a list, more like comma separated variables
    A point, x,y

    Returns
    -----
    first_derivative: np.array
    The first derivative (gradient) at the point entered
    """
    x, y = point
    dfdx = -2 * (1-x) - 40 * x * (y-x**2)
    dfdy = 20 * (y-x**2)

    return np.array([dfdx, dfdy])

```

3a

```

In [16]: steepest_descent(rosenbrock_function3d, rosenbrock_gradient, np.array([-0.5,
func:'steepest_descent' took: 0.0093 sec
Out[16]: {'Starting point': array([0.99999089, 0.99998153]),
'Evaluation': 8.361266796946337e-11,
'Path to minimum': array([[ -0.5          ,  1.5          ],
[ -1.05         ,  0.875         ],
[ -0.845175     ,  0.94325      ],
...,
[ 0.99999068    ,  0.99998135   ],
[ 0.99999093    ,  0.99998135   ],
[ 0.99999089    ,  0.99998153   ]]),
'Steps to converge': 1205}

```

3b

```

In [17]: @timeit

def stochastic_gradient_descent(func, func_gradient, x0, alpha=0.1, tol=1e-5, steps=1000):
    """
    This function finds a local minimum using the stochastic gradient method.
    Parameters
    -----
    func:
    The function whose minima we plan to find (inputted as a function)

    func_gradient:

```

first_derivative of func

x0: np.array

position from which we start searching for the minima

alpha: float

step size

tol: float

tolerance value

stochastic_injection: 0 or 1

controls the magnitude of stochasticity (multiplied with stochastic_deriv)
0 for no stochasticity, equivalent to SD.

Returns

Starting point: entered x0

Evaluation: value of function at x0

Path to minimum: lists all points evaluated on the way to minimum

Steps to converge: counts number of steps until local minimum is reached
...

evaluate the gradient at starting point

deriv = func_gradient(x0)

count=0

visited=[x0]

while LA.norm(deriv) > tol and count < 1e5:

 if stochastic_injection>0:

formulate a stochastic_deriv that is the same norm as your gradient

 dim = deriv.shape

 stochastic_deriv = np.random.random(dim) * 2 - 1

 stochastic_norm = LA.norm(stochastic_deriv)

 stochastic_deriv = stochastic_deriv / stochastic_norm * LA.norm(deriv)

 else:

 stochastic_deriv = np.zeros(len(x0))

 direction = -(deriv + stochastic_injection * stochastic_deriv)

calculate new point position

 x1 = x0 - direction * alpha

 if func(x1) < func(x0):

Check if new value is less than previous value. If so, this is

and start the next search step from the current value

 alpha = alpha * 1.2

 x0 = x1

 visited.append(x1)

 deriv = func_gradient(x1)

 print(f'good step {x1}')

 else:

If new_value > prev_value, we are moving in the wrong direction

and redo that step starting from the previous value.

 alpha = alpha * 0.5

 print(f'bad step {x1}')

```

        count+=1
    return {"x":x0,"evaluation":func(x0),"path":np.asarray(visited), "Number

```

```
In [18]: stochastic_gradient_descent(rosenbrock_function3d, rosenbrock_gradient, np.a
```

```
func:'stochastic_gradient_descent' took: 0.0183 sec
```

```

Out[18]: {'x': array([0.99999089, 0.99998153]),
  'evaluation': 8.361266796946337e-11,
  'path': array([[ -0.5         ,  1.5         ],
    [-1.05         ,  0.875         ],
    [-0.845175      ,  0.94325        ],
    ...,
    [ 0.99999068,  0.99998135],
    [ 0.99999093,  0.99998135],
    [ 0.99999089,  0.99998153]]),
  'Number of steps': 1205}

```

3c

```
In [19]: x0 = np.array([-0.5, 1.5])
```

```

CG_rosenbrock = minimize(rosenbrock_function3d, x0, method='CG', options={'g
print(CG_rosenbrock)

```

```
Optimization terminated successfully.
```

```
Current function value: 0.000000
```

```
Iterations: 20
```

```
Function evaluations: 132
```

```
Gradient evaluations: 44
```

```
message: Optimization terminated successfully.
```

```
success: True
```

```
status: 0
```

```
fun: 2.0711221375743512e-13
```

```
x: [ 1.000e+00  1.000e+00]
```

```
nit: 20
```

```
jac: [ 4.992e-08 -2.474e-08]
```

```
nfev: 132
```

```
njev: 44
```

```

In [20]: BFGS_rosenbrock = minimize(rosenbrock_function3d, x0, method='BFGS', options
print(BFGS_rosenbrock)

```

```

Optimization terminated successfully.
    Current function value: 0.000000
    Iterations: 22
    Function evaluations: 93
    Gradient evaluations: 31
message: Optimization terminated successfully.
success: True
status: 0
    fun: 1.6857105436734322e-13
     x: [ 1.000e+00  1.000e+00]
    nit: 22
    jac: [ 1.153e-07 -1.294e-08]
hess_inv: [[ 5.099e-01  1.020e+00]
            [ 1.020e+00  2.089e+00]]
    nfev: 93
    njev: 31

```

In the case of the Rosenbrock Banana Function, CG and BFGS are far better than SGD. SGD found the minimum in ~1200 steps, while CG and BFGS took 132 and 93, respectively.

3d

No. Because of the stochasticity, number of steps with SGD will vary largely and so you need to take an average value after a few runs for comparison.

3e

I found the performance of SGD to be more consistent than that of steepest descent. For example, at point (-1.5, 1.5), SGD and steepest descent took ~1200 steps to converge. However, as the numbers got bigger the difference in step size grew. At point (-150, 150), SGD took ~79000 steps to converge while steepest descent took ~490000. At (-1500, 1500), SGD still took ~79000 steps while steepest descent took ~790000.

However when comparing CD and BFGS with SGD, their performances (in terms of step sizes) did not increase as quickly as the step size of steepest descent did.

Question 4: Stochastic Gradient Descent with Momentum (SGDM)

```

In [21]: def momentum_function3d(point):
         """
         This function calculates the value of a function of x and y at a certain
         Parameters
         -----
         x, y: the point to evaluate the function at

         Return
         -----
         the function value at the point

```

```

"""
x, y = point

return 2*x**2 - 1.05*x**4 + x**6/6 + x*y + y**2

```

```

In [22]: def momentum_gradient(point):
        """
        This function calculates the first derivative of a function at a given point.
        Parameters
        -----
        point: list. Not really a list, more like comma separated variables
        A point, x,y

        Returns
        -----
        first_derivative: np.array
        The first derivative (gradient) at the point entered
        """
        x, y = point
        dfdx = 4*x - 4.2*x**3 + x**5 + y
        dfdy = x + 2*y

        return np.array([dfdx, dfdy])

```

4a

```

In [23]: stochastic_gradient_descent(momentum_function3d, momentum_gradient, np.array([0.5, 0.5]))
func:'stochastic_gradient_descent' took: 0.0020 sec

```

```

Out[23]: {'x': array([-1.74755328,  0.87377865]),
          'evaluation': 0.29863844224600855,
          'path': array([[ -1.5       ,  1.5       ],
                        [-1.708125 ,  1.35       ],
                        [-1.81711465,  1.230975  ],
                        [-1.72366291,  1.13811871],
                        [-1.81648029,  1.04263383],
                        [-1.73077839,  1.01476596],
                        [-1.77260058,  0.97759624],
                        [-1.73965321,  0.95033542],
                        [-1.77022043,  0.92148765],
                        [-1.74397071,  0.91366684],
                        [-1.75467953,  0.90291347],
                        [-1.74553402,  0.89499619],
                        [-1.7540083 ,  0.88673796],
                        [-1.74656205,  0.88456827],
                        [-1.74961089,  0.88154912],
                        [-1.74682193,  0.87938454],
                        [-1.74960838,  0.87708366],
                        [-1.74708789,  0.87655687],
                        [-1.74825522,  0.8757213 ],
                        [-1.74715493,  0.87519094],
                        [-1.74777788,  0.87486877],
                        [-1.74758021,  0.87463399],
                        [-1.74765485,  0.87439135],
                        [-1.74754601,  0.87419677],
                        [-1.74764902,  0.87402131],
                        [-1.74753349,  0.87397242],
                        [-1.74759689,  0.87391111],
                        [-1.74752418,  0.8738708 ],
                        [-1.74757107,  0.87384747],
                        [-1.74755094,  0.87383152],
                        [-1.74756213,  0.87381419],
                        [-1.74754713,  0.87380191],
                        [-1.74755705,  0.8737956 ],
                        [-1.74755201,  0.87379104],
                        [-1.74755504,  0.87378622],
                        [-1.74755067,  0.87378288],
                        [-1.74755379,  0.87378114],
                        [-1.74755205,  0.87377996],
                        [-1.74755328,  0.87377865]]),
          'Number of steps': 39}

```

```

In [24]: x0 = np.array([-1.5, -1.5])

CG_momentum = minimize(momentum_function3d, x0, method='CG', options={'gtol'
print(CG_momentum)

```

```

Optimization terminated successfully.
    Current function value: 0.298638
    Iterations: 7
    Function evaluations: 63
    Gradient evaluations: 21
message: Optimization terminated successfully.
success: True
status: 0
    fun: 0.29863844223965763
     x: [-1.748e+00  8.738e-01]
    nit: 7
   jac: [ 8.404e-06  7.227e-07]
  nfev: 63
  njev: 21

```

```
In [25]: BFGS_momentum = minimize(momentum_function3d, x0, method='BFGS', options={'c
print(BFGS_momentum)
```

```

Optimization terminated successfully.
    Current function value: 0.298638
    Iterations: 8
    Function evaluations: 30
    Gradient evaluations: 10
message: Optimization terminated successfully.
success: True
status: 0
    fun: 0.29863844223686065
     x: [-1.748e+00  8.738e-01]
    nit: 8
   jac: [ 1.341e-07 -7.451e-09]
 hess_inv: [[ 8.569e-02 -4.290e-02]
            [-4.290e-02  5.109e-01]]
  nfev: 30
  njev: 10

```

On average, stochastic gradient descent did better than conjugate gradients (~40 steps VS ~60 steps). Stochastic gradient and BFGS had roughly the same performance (~40 steps and ~30 steps)

4b

```
In [26]: @timeit

def SGDM(func, func_gradient, x0, alpha=0.1, gamma=0.9, tol=1e-5, stochastic_injec
    ...
    This function finds a local minimum using the stochastic gradient method
    Parameters
    -----
    func:
    The function whose minima we plan to find (inputted as a function)

    func_gradient:
    first_derivative of func

    x0: np.array
```

position from which we start searching for the minima

alpha: float
step size

gamma: float
momentum value

tol: float
tolerance value

stochastic_injection: 0 or 1
controls the magnitude of stochasticity (multiplied with stochastic_deriv)
0 for no stochasticity, equivalent to SD.

Returns

Starting point: entered x0
Evaluation: value of function at x0
Path to minimum: lists all points evaluated on the way to minimum
Steps to converge: counts number of steps until local minimum is reached
...

evaluate the gradient at starting point
deriv = func_gradient(x0)

count=0

visited=[x0]

while LA.norm(deriv) > tol and count < 1e5:

 if stochastic_injection>0:

formulate a stochastic_deriv that is the same norm as your gradient
 stochastic_deriv= np.random.random(2) * 2 - 1
 stochastic_norm = LA.norm(stochastic_deriv)
 stochastic_deriv = stochastic_deriv / stochastic_norm * LA.norm(deriv)

 else:

 stochastic_deriv=np.zeros(len(x0))

 if count == 0:

 previous_direction = -deriv

 direction=-(deriv+stochastic_injection*stochastic_deriv + gamma * previous_direction)
 x1 = x0 + alpha * direction

 if func(x1) < func(x0):

Check if new value is less than previous value. If so, this is a new minimum
 # and start the next search step from the current value

 alpha = alpha * 1.2

 x0 = x1

 visited.append(x1)

 deriv = func_gradient(x0)

 else:

If new_value > prev_value, we are moving in the wrong direction
 # and redo that step starting from the previous value.

 if alpha<1e-5:

 previous_direction=previous_direction-previous_direction

 else:

do the same as SGD here

 alpha = alpha * 0.5


```
        count+=1  
    return {"x":x0,"Evaluation":func(x0),"Path":np.asarray(visited), "Number
```

```
In [27]: SGDM(momentum_function3d, momentum_gradient, np.array([-1.5, -1.5]))
```

func:'SGDM' took: 0.0030 sec

```

Out[27]: {'x': array([-1.74755242,  0.87377317]),
          'Evaluation': 0.2986384422461035,
          'Path': array([[ -1.5          , -1.5          ],
                        [-1.42887291, -0.99991264],
                        [-1.52341299, -0.95558559],
                        [-1.51358025, -0.93314037],
                        [-1.50819735, -0.90514635],
                        [-1.51171835, -0.87149957],
                        [-1.51193333, -0.8702979 ],
                        [-1.5138259 , -0.86970886],
                        [-1.51209007, -0.86936447],
                        [-1.51013788, -0.86869735],
                        [-1.50914783, -0.86659327],
                        [-1.50780145, -0.86603561],
                        [-1.50810051, -0.86426464],
                        [-1.50881599, -0.86328878],
                        [-1.50874209, -0.86205446],
                        [-1.50871703, -0.86196794],
                        [-1.50859369, -0.86197152],
                        [-1.50861021, -0.86193535],
                        [-1.50862084, -0.86188384],
                        [-1.50860418, -0.86182104],
                        [-1.50856939, -0.86180588],
                        [-1.50857772, -0.86180438],
                        [-1.50862888, -0.86175653],
                        [-1.50857728, -0.86163747],
                        [-1.50853016, -0.8614841 ],
                        [-1.50859024, -0.86131583],
                        [-1.5086977 , -0.86120214],
                        [-1.50868289, -0.86091939],
                        [-1.50877758, -0.8606204 ],
                        [-1.50864942, -0.86058901],
                        [-1.50865986, -0.86010272],
                        [-1.5083494 , -0.85988323],
                        [-1.50825608, -0.85918581],
                        [-1.50857053, -0.85853138],
                        [-1.50877006, -0.85847055],
                        [-1.50849573, -0.85843156],
                        [-1.50899962, -0.85724743],
                        [-1.50935077, -0.85561289],
                        [-1.51029808, -0.8547077 ],
                        [-1.51118549, -0.85268783],
                        [-1.50958643, -0.85082981],
                        [-1.50861662, -0.8506436 ],
                        [-1.50822936, -0.85064252],
                        [-1.51060204, -0.84818799],
                        [-1.513403  , -0.84562414],
                        [-1.50988777, -0.83994568],
                        [-1.50739991, -0.8315602 ],
                        [-1.51145361, -0.82912219],
                        [-1.51563708, -0.82707445],
                        [-1.51740927, -0.82671246],
                        [-1.52416621, -0.81224243],
                        [-1.52789236, -0.79137261],
                        [-1.5139853 , -0.77658175],
                        [-1.50102341, -0.7515267 ]],
          )

```

```

[-1.50348395, -0.75121575],
[-1.51786004, -0.74459859],
[-1.54127613, -0.72579468],
[-1.52719251, -0.66603555],
[-1.51817571, -0.66527007],
[-1.53778436, -0.58585119],
[-1.56126312, -0.49547162],
[-1.52436188, -0.39937429],
[-1.47632375, -0.36966609],
[-1.50214622, -0.23476685],
[-1.44060314, -0.16033728],
[-1.4122792 , -0.14722699],
[-1.37314195,  0.0151895 ],
[-1.4798359 ,  0.15702739],
[-1.6047083 ,  0.33006197],
[-1.73508535,  0.49289655],
[-1.77033475,  0.63790684],
[-1.64515048,  0.78308844],
[-1.76231657,  0.83695996],
[-1.74609625,  0.86784473],
[-1.74631626,  0.87044695],
[-1.7476742 ,  0.87151245],
[-1.74729029,  0.87140865],
[-1.74700886,  0.87212268],
[-1.74809577,  0.87290828],
[-1.74702504,  0.87323309],
[-1.74733609,  0.87281383],
[-1.74728185,  0.87308691],
[-1.7474133 ,  0.87351685],
[-1.74759802,  0.87362149],
[-1.74753497,  0.87372373],
[-1.74754153,  0.87373789],
[-1.74754202,  0.87374692],
[-1.74754606,  0.87374306],
[-1.74754335,  0.87374985],
[-1.74755902,  0.8737628 ],
[-1.74755931,  0.87377075],
[-1.74754841,  0.87376586],
[-1.74755381,  0.87377095],
[-1.74755244,  0.87377015],
[-1.74755277,  0.87377151],
[-1.74755084,  0.87377256],
[-1.74755109,  0.87377346],
[-1.74755242,  0.87377317]]),
'Number of step to converge': 98}

```

4b

No, I did not get a better result when using SGDM. When comparing number of steps, SGD took the fewest number of steps to find global minimum. This was unexpected, as I thought that SGD with momentum would perform better. Since the momentum takes previous good steps into consideration, I thought momentum would serve as a guiding force to move the search in the right direction.