

```
In [38]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
from sklearn.preprocessing import StandardScaler
import random
import math
import warnings

warnings.filterwarnings("ignore")

%matplotlib inline
```

1a

```
In [39]: import sklearn.model_selection as skl_model

admit_predict = pd.read_csv('Admission_Predict_Ver1.1.csv')

admit_predict.set_index('Serial No.', inplace=True)
admit_predict.head()

# LOR = letter of rec
# SOP = statement of purpose
```

```
Out[39]:
```

	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
Serial No.								
1	337	118	4	4.5	4.5	9.65	1	0.92
2	324	107	4	4.0	4.5	8.87	1	0.76
3	316	104	3	3.0	3.5	8.00	1	0.72
4	322	110	3	3.5	2.5	8.67	1	0.80
5	314	103	2	2.0	3.0	8.21	0	0.65

Features related to chance of admission are GRE score, TOEFL score, university rating, SOP, LOR, CGPA and Research.

```
In [26]: rankings = admit_predict['Chance of Admit'].values

features = admit_predict.drop(columns='Chance of Admit').values
scaler = StandardScaler()
features_norm = scaler.fit_transform(features)
```

1b

```
In [27]: def tanh(x):
          return np.tanh(x)

def tanh_grad(x):
    return 1-np.tanh(x)**2
```

```
In [40]: class simple_perceptron():
          def __init__(self, input_dim, output_dim, learning_rate=0.01, activation='tanh'):

              self.input_dim = input_dim
              self.output_dim = output_dim
              self.activation = activation
              # "The activation gradient is a measure of how sensitive the activation is to changes in the input"
              self.activation_grad = activation_grad
              self.lr = learning_rate
              ### initialize parameters ###
              self.weights = np.random.uniform(0, 1, (self.input_dim, self.output_dim))
              self.biases = np.random.uniform(0, 0.05, self.output_dim)

          def predict(self, X):
              if len(X.shape) == 1:
                  X = X.reshape((-1,1))
              dim = X.shape[1]
              # Check that the dimension of accepted input data is the same as expected
              if not dim == self.input_dim:
                  raise Exception("Expected input size %d, accepted %d!"%(self.input_dim, dim))
              ### Calculate logit and activation ###
              self.z = (X @ self.weights) + self.biases #shape(X.shape[0],1)
              self.a = self.activation(self.z) #shape(X.shape[0],1)
              return self.a

          def fit(self, X, y):
              # Transform the single-sample data into 2-dimensional, for the convolution
              if len(X.shape) == 1:
                  X = X.reshape((-1,1))
              if len(y.shape) == 1:
                  y = y.reshape((-1,1))
              self.predict(X)
              # subtracts true y values from predicted values (self.a)
              errors = (self.a - y) * self.activation_grad(self.z)
              # matrix multiplication between transpose of errors and input data X
              weights_grad = errors.T.dot(X)
              # sums up errors along rows (axis 0)
              bias_grad = np.sum(errors, axis = 0)
              ### Update weights and biases from the gradient ###
              self.weights -= self.lr * weights_grad.T
              self.biases -= self.lr * bias_grad.T

          def train_on_epoch(self, X, y, batch_size=32):
              # Every time select batch_size samples from the training set, until the end of the set
              order = list(range(X.shape[0]))
              random.shuffle(order)
              n = 0
              while n < math.ceil(len(order)/batch_size)-1: # Parts that can fill the batch
```

```

        self.fit(X[order[n*batch_size:(n+1)*batch_size]],y[order[n*batch
        n += 1
    # Parts that cannot fill one batch
    self.fit(X[order[n*batch_size:]],y[order[n*batch_size:]])

def evaluate(self,X,y):
    # Transform the single-sample data into 2-dimensional
    if len(X.shape) == 1:
        X = X.reshape((1,-1))
    if len(y.shape) == 1:
        y = y.reshape((1,-1))
    ### means square error ###
    return np.mean((self.predict(X)-y)**2)

def get_weights(self):
    return (self.weights,self.biases)

def set_weights(self,weights):
    self.weights = np.array(weights[0])
    self.biases = np.array(weights[1])

```

1c

```

In [45]: from sklearn.model_selection import train_test_split,KFold

def Kfold(k, Xs, ys, epochs, learning_rate = 0.0001, draw_curve=True):
    # The total number of examples for training the network
    total_num = len(Xs)

    # Built in K-fold function in Sci-Kit Learn
    kf = KFold(n_splits = k, shuffle=True)
    # record error for each model
    train_error_all = []
    test_error_all = []

    for train_selector,test_selector in kf.split(range(total_num)):
        ### Decide training examples and testing examples for this fold ###
        train_Xs = Xs[train_selector]
        test_Xs = Xs[test_selector]
        train_ys = ys[train_selector]
        test_ys = ys[test_selector]

        val_array = []
        # Split training examples further into training and validation
        train_in, val_in, train_real, val_real = train_test_split(train_Xs,t

        ### Establish the model for simple perceptron here ###
        model = simple_perceptron(train_Xs.shape[1], 1)

        # Save the lowest weights, so that we can recover the best model
        weights = model.get_weights()
        lowest_val_err = np.inf
        for _ in range(epochs):
            # Train model on a number of epochs, and test performance in the

```

```

        model.train_on_epoch(train_in,train_real)
        val_err = model.evaluate(val_in,val_real)
        val_array.append(val_err)
        if val_err < lowest_val_err:
            lowest_val_err = val_err
            weights = model.get_weights()

# The final number of epochs is when the minimum error in validation
final_epochs = val_array.index(min(val_array)) + 1 # +1 to consider
print("Number of epochs with lowest validation:",final_epochs)
# Recover the model weight
model.set_weights(weights)

# Report result for this fold
train_error = model.evaluate(train_Xs, train_ys)
train_error_all.append(train_error)
test_error = model.evaluate(test_Xs, test_ys)
test_error_all.append(test_error)
print("Train error:",train_error)
print("Test error:",test_error)

if draw_curve:
    plt.figure()
    plt.plot(np.arange(len(val_array))+1,val_array,label='Validation')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.legend()

print("Final results:")
print("Training error:%f+-%f"%(np.average(train_error_all),np.std(train_
print("Testing error:%f+-%f"%(np.average(test_error_all),np.std(test_err

def show_correlation(xs,ys):
    xs = model.predict(xs).reshape(-1,)
    ys = ys.reshape(-1)
    plt.figure()
    plt.scatter(xs,ys,s=0.5)
    r = [np.min([np.min(xs),np.min(ys)]),np.max([np.max(xs),np.max(ys)])]
    plt.plot(r,r,'r')
    plt.xlabel("Predictions")
    plt.ylabel("Ground truth")
    corr=np.corrcoef([xs,ys])[1,0]
    print("Correlation coefficient:",corr)

# return the last model
return model

```

```


In [46]: def tanh(x):
        return np.tanh(x)



def tanh_grad(x):
    return 1-np.tanh(x)**2



model = Kfold(5, features_norm, rankings, 50, learning_rate = 0.00001, draw_
print(model)

```

Number of epochs with lowest validation: 21
Train error: 0.03420285585867554
Test error: 0.02833616324911934
Number of epochs with lowest validation: 14
Train error: 0.03212957923900913
Test error: 0.03823247157429192
Number of epochs with lowest validation: 29
Train error: 0.03396143696543532
Test error: 0.03726124318468313
Number of epochs with lowest validation: 40
Train error: 0.03360303355729738
Test error: 0.03633624576843458
Number of epochs with lowest validation: 29
Train error: 0.035760189004682995
Test error: 0.027995262080891547
Final results:
Training error: 0.033931+-0.001164
Testing error: 0.033632+-0.004505
<__main__.simple_perceptron object at 0x13a7e6290>








```
In [47]: show_correlation(features_norm, rankings)
```

Correlation coefficient: 0.9057438961067454

```
Out[47]: <__main__.simple_perceptron at 0x13a7e6290>
```







Yes, I think the features are a good indicator of getting into graduate school. Now we remove GRE Scores and test again.


```
In [50]: features_no_GRE = admit_predict.drop(columns=['Chance of Admit', 'GRE Score'])
scaler = StandardScaler()
features_norm_no_GRE = scaler.fit_transform(features)

model2 = Kfold(5, features_norm_no_GRE, rankings, 50, learning_rate = 0.00001)
print(model2)
```

```
Number of epochs with lowest validation: 50
Train error: 0.029522150495362807
Test error: 0.02922318051672919
Number of epochs with lowest validation: 15
Train error: 0.044307049194561035
Test error: 0.0427414446593212
Number of epochs with lowest validation: 43
Train error: 0.03766941213285055
Test error: 0.034889763959733484
Number of epochs with lowest validation: 6
Train error: 0.03574767343904443
Test error: 0.039825314774456444
Number of epochs with lowest validation: 16
Train error: 0.03616130897572605
Test error: 0.03457722323799468
Final results:
Training error:0.036682+-0.004723
Testing error:0.036251+-0.004668
<__main__.simple_perceptron object at 0x13a28f5d0>
```






```
In [51]: show_correlation(features_norm_no_GRE, rankings)
```

Correlation coefficient: 0.9057438961067454

```
Out[51]: <__main__.simple_perceptron at 0x13a7e6290>
```



I guess GRE Scores are not that important; the correlation graphs look more or less identical

2a

This dataset has a lot more categorical data than the previous one, which just contained numbers and binary values. There were more steps to prepare and separate the features and rankings.

```
In [52]: titanic = pd.read_csv('titanic.csv')

titanic.set_index('PassengerId', inplace=True)
titanic = titanic.dropna()

titanic_numericals = titanic.drop(['Ticket', 'Name', 'Sex'], axis=1)
```

```
In [76]: from sklearn.preprocessing import OneHotEncoder

# normalize numerical data
normalized_numerical = titanic_numericals.iloc[:, 1:6].apply(lambda x: (x -

# one hot encode categorical data
encoder = OneHotEncoder(handle_unknown='ignore')
encoded_categoricals = encoder.fit_transform(titanic[['Sex', 'Embarked']]).toarray()
encoded_categoricals_df = pd.DataFrame(encoded_categoricals)

titanic_transformed = pd.concat([normalized_numerical, encoded_categoricals_df], axis=1)
```

```
titanic_transformed = titanic_transformed.dropna()  
  
#titanic_transformed.head()
```



2b



Use the simple perceptron model we developed in Q1. Use 80% of the data for training and 20% of the data for testing and do 5-fold validation. Can we predict who will survive? Play around with the features to determine which ones give you a better chance to get back to shore


```
In [83]: features = titanic_transformed.values  
         predictions = titanic['Survived'].values  
  
         Kfold(5, features, predictions, 60)
```

```
Number of epochs with lowest validation: 25  
Train error: 0.39564764868540236  
Test error: 0.31329868217727747  
Number of epochs with lowest validation: 27  
Train error: 0.3679870596334657  
Test error: 0.4782723493771518  
Number of epochs with lowest validation: 14  
Train error: 0.34261610818235694  
Test error: 0.40184007160530455  
Number of epochs with lowest validation: 12  
Train error: 0.3887770023802423  
Test error: 0.366983306429162  
Number of epochs with lowest validation: 3  
Train error: 0.37021299484898945  
Test error: 0.35849825710336575  
Final results:  
Training error:0.373048+-0.018536  
Testing error:0.383779+-0.055034
```

```
Out[83]: <__main__.simple_perceptron at 0x13a4bab90>
```







```
In [ ]: show_correlation(features, predictions)
```

Yes, we can fairly accurately predict who will survive. Although I could not get my correlation graph to work, I say this based on the loss function. The largest determinant for survival was Sex, as it had the largest weight.

3a


```
In [64]: import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D

def generate_X(number):
    xs = (np.random.random(number)*2-1) * 10
    return xs

def generate_data(number, stochascity=0.05):
    xs = generate_X(number)
    fs = 3 * np.sin(xs)-5
    stochastic_ratio = (np.random.random(number)*2-1) * stochascity+1
    return xs, fs * stochastic_ratio
```

```
In [65]: x, y = generate_data(5000, 0.1)
plt.scatter(x, y, s=0.1)
```

```
Out[65]: <matplotlib.collections.PathCollection at 0x13a4d2810>
```

In [84]: `x, y = generate_data = (1000, 0.1)`

3b

```
In [66]: #from sklearn.neural_network import MLPRegressor
#from sklearn.model_selection import train_test_split, KFold

import numpy as np
from sklearn.neural_network import MLPRegressor
from sklearn.model_selection import cross_val_score
from sklearn.metrics import mean_squared_error

x, y = generate_data(1000, 0.1)

MLP = MLPRegressor(hidden_layer_sizes=(8,), random_state = 49)

mean_squared_error = -cross_val_score(MLP, x.reshape(-1, 1), y, cv = 5, scor


for fold, mse in enumerate(mean_squared_error):
    print(f'Fold {fold + 1}: MSE = {mse}')

MLP.fit(x.reshape(-1, 1), y)
y_pred = MLP.predict(x.reshape(-1, 1))

plt.figure()
plt.scatter(x, y, label='Observed')
plt.scatter(x, y_pred, label='Model Prediction', alpha=0.6)
plt.xlabel('x values')
plt.ylabel('y values')
```

```
plt.legend()
plt.show()
```

Fold 1: MSE = 6.557191509323998
 Fold 2: MSE = 5.759326237994227
 Fold 3: MSE = 6.967023793409838
 Fold 4: MSE = 6.604969312598687
 Fold 5: MSE = 6.710412821254925



3c

The model prediction is not good at all. Let's add more hidden layers to see if that improved our predictions.

```
In [67]: # second try with more hidden layers

import numpy as np
from sklearn.neural_network import MLPRegressor
from sklearn.model_selection import cross_val_score
from sklearn.metrics import mean_squared_error

x, y = generate_data(1000, 0.1)

MLP = MLPRegressor(hidden_layer_sizes=(1000,), random_state = 49)


mean_squared_error = -cross_val_score(MLP, x.reshape(-1, 1), y, cv = 5, scoring='neg_mean_squared_error')

for fold, mse in enumerate(mean_squared_error):
    print(f'Fold {fold + 1}: MSE = {mse}')
```

```
MLP.fit(x.reshape(-1, 1), y)
y_pred = MLP.predict(x.reshape(-1, 1))

plt.figure()
plt.scatter(x, y, label='Observed')
plt.scatter(x, y_pred, label='Model Prediction', alpha=0.6)
plt.xlabel('X')
plt.ylabel('y')
plt.legend()
plt.show()
```

Fold 1: MSE = 3.0936014218491175
Fold 2: MSE = 2.6880131156572475
Fold 3: MSE = 3.0773700882193693
Fold 4: MSE = 3.228400144248916
Fold 5: MSE = 2.4915781161043653



I increased the number of hidden layers to 1000 and this greatly improved the model. However, it's still not a good fit to the data despite having very many hidden layers (1000). In real life, I don't know if this is even a realistic number of hidden layers to have in a model.